

PH125.9x - Capstone: MovieLens Project

Ian Espejo-Campos

Abstract

This project creates a movie recommendation system using the 10M version of the MovieLens dataset. In addition to movie and user effects, 'title' and 'timestamp' variables are used to extract movies' release years, and calculate a new rate timestamp indicator, incorporating all four effects in the final model, which are also regularized by a tuning parameter of 0.5 chosen through cross-validation. The final model reaches a residual mean squared error (RMSE) ~ 0.8647 when employed on the Validation Set, which is lower than the project's target.

Keywords: movielens, recommendation system, EDA, capstone, r markdown

1 Introduction

MovieLens (<https://movielens.org/>) is a website run by GroupLens, a research lab at the University of Minnesota. It works as a collaborative movie recommendation system, based on their film preferences or taste profile, measured employing Machine Learning tools on users' movie ratings and movie reviews (Harper & Konstan, 2015).

In this project, we will be working with the MovieLens 10M dataset:

This data set contains 10000054 ratings and 95580 tags applied to 10681 movies by 71567 users of the online movie recommender service MovieLens.

Users were selected at random for inclusion. All users selected had rated at least 20 movies. Unlike previous MovieLens data sets, no demographic information is included. Each user is represented by an id, and no other information is provided.

The data are contained in three files, movies.dat, ratings.dat and tags.dat. Also included are scripts for generating subsets of the data to support five-fold cross-validation of rating predictions. More details about the contents and use of all these files follows.

The project approach is to build a model using a train set that minimizes the loss function by considering the following effects:

- *Movie*: The rating can be predicted based on other users' evaluation of the movie.
- *User*: It can also be predicted by the user's usual rating to watched movies.
- *Release Year*: Improves the algorithm by grouping movies with its yearly competitors.
- *Rate Timestamp*: Number of years between release and moment when the rate was given, improving the model by considering the movie's gained popularity.

After that, the model is regularized with a tuning parameter chosen through cross-validation, in order to penalize large estimates that are formed using small sample sizes (Irizarry, 2020). A residual mean squared error (RMSE) was used as the loss function (the typical error we make when predicting a movie rating) for performance measure.

2 Methodology

As the first step, we downloaded the MovieLens dataset with the code provided in the course. Once we had a train set, we analyzed some of its properties through descriptive statistics and histograms, in order to gain insights for the prediction model. Then, we added the movie's release year and rate timestamp to the dataset, in order to use them as predictors.

Finally, the model was built using movie, user, release year and rate timestamp effects on the edx set. The effects were later regularized looking for a tuning parameter (λ) that minimized the RMSE. Every intermediate result obtained was compared to the $RMSE < 0.8649$ goal.

2.1 Data extraction

Code provided on the Course/Capstone Project: All Learners/Project Overview: MovieLens section of the PH125.9x - Capstone: MovieLens Project course.

```
# Import libraries
rm(list=ls())

if(!require(tidyverse))
  install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret))
  install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table))
  install.packages("data.table", repos = "http://cran.us.r-project.org")
if(!require(ggplot2))
  install.packages("ggplot2", repos = "http://cran.us.r-project.org")
if(!require(knitr))
  install.packages("knitr", repos = "http://cran.us.r-project.org")

library(tidyverse)
library(caret)
```

```

library(data.table)
library(ggplot2)
library(knitr)

# Get from URL or download file in Documents
zip_file <- "~/ml-10m.zip"
if(file.exists(zip_file)) {dl <- zip_file} else {
  dl <- tempfile()
  download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip",
    dl)
}

# Unzip and read data
ratings <- fread(text = gsub(":", "\t",
                           readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
  col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl,
                                         "ml-10M100K/movies.dat")), "\\:", 3)
colnames(movies) <- c("movieId", "title", "genres")

movies <- as.data.frame(movies) %>%
  mutate(movieId = as.numeric(movieId),
    title = as.character(title),
    genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind = "Rounding")
test_index <- createDataPartition(y = movielens$rating,
  times = 1,
  p = 0.1,
  list = F)

edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

```

2.2 Data description

The edx train set contains 9,000,055 rows and 6 columns, each variable meaning: userId of the rater, movieId, title and genres of the movie rated, timestamp of the information, and the rating gave by the user, which goes from 0.5 to 5.0 stars.

```
dim(edx)
```

```
## [1] 9000055      6
```

```
head(edx) %>% as.tibble()
```

```
## # A tibble: 6 x 6
##   userId movieId rating timestamp title          genres
##   <int>   <dbl>   <dbl>     <int> <chr>        <chr>
## 1     1     122     5 838985046 Boomerang (1992) Comedy|Romance
## 2     1     185     5 838983525 Net, The (1995) Action|Crime|Thriller
## 3     1     292     5 838983421 Outbreak (1995) Action|Drama|Sci-Fi|T~
## 4     1     316     5 838983392 Stargate (1994) Action|Adventure|Sci-~
## 5     1     329     5 838983392 Star Trek: Generations~ Action|Adventure|Dram~
## 6     1     355     5 838984474 Flintstones, The (1994) Children|Comedy|Fanta~
```

We can analyze the train set using some descriptive statistics.

```
summary(edx)
```

```
##      userId      movieId      rating      timestamp
## Min.   :    1   Min.   :    1   Min.   :0.500   Min.   :7.897e+08
## 1st Qu.:18124  1st Qu.:   648  1st Qu.:3.000   1st Qu.:9.468e+08
## Median :35738  Median :  1834  Median :4.000   Median :1.035e+09
## Mean   :35870  Mean   :   4122  Mean   :3.512   Mean   :1.033e+09
## 3rd Qu.:53607  3rd Qu.:  3626  3rd Qu.:4.000   3rd Qu.:1.127e+09
## Max.    :71567  Max.    :65133  Max.    :5.000   Max.    :1.231e+09
##      title      genres
## Length:9000055   Length:9000055
## Class :character Class :character
## Mode  :character Mode  :character
##
##
##
```

Notice that there are 69,878 and 10,677 different users and movies, respectively.

```
summarize(edx,
  unique_users = n_distinct(userId),
  unique_movies = n_distinct(movieId))
```

```
##   unique_users unique_movies
## 1         69878         10677
```

Here we can see the top 50 best ranking movies, and how many votes each one received.

```
edx %>% group_by(movieId) %>%
  summarise(stars = mean(rating),
    votes = n()) %>%
  left_join(edx, by = "movieId") %>%
  select(movieId, title, votes, stars) %>%
  unique() %>%
  filter(title != "NA") %>%
  arrange(desc(stars)) %>%
  slice(1:50) %>%
  mutate(ranking = 1:n()) %>%
  select(ranking, everything()) %>%
  kable()
```

| ranking | movieId | title | votes | stars |
|---------|---------|---|-------|----------|
| 1 | 3226 | Hellhounds on My Trail (1999) | 1 | 5.000000 |
| 2 | 33264 | Satan's Tango (S  t  ntang   ) (1994) | 2 | 5.000000 |
| 3 | 42783 | Shadows of Forgotten Ancestors (1964) | 1 | 5.000000 |
| 4 | 51209 | Fighting Elegy (Kenka erejii) (1966) | 1 | 5.000000 |
| 5 | 53355 | Sun Alley (Sonnenallee) (1999) | 1 | 5.000000 |
| 6 | 64275 | Blue Light, The (Das Blaue Licht) (1932) | 1 | 5.000000 |
| 7 | 5194 | Who's Singin' Over There? (a.k.a. Who Sings Over There) (Ko to tamo peva) (1980) | 4 | 4.750000 |
| 8 | 26048 | Human Condition II, The (Ningen no joken II) (1959) | 4 | 4.750000 |
| 9 | 26073 | Human Condition III, The (Ningen no joken III) (1961) | 4 | 4.750000 |
| 10 | 65001 | Constantine's Sword (2007) | 2 | 4.750000 |
| 11 | 4454 | More (1998) | 7 | 4.714286 |
| 12 | 5849 | I'm Starting From Three (Ricomincio da Tre) (1981) | 3 | 4.666667 |
| 13 | 63808 | Class, The (Entre les Murs) (2008) | 3 | 4.666667 |
| 14 | 7452 | Mickey (2003) | 1 | 4.500000 |
| 15 | 7823 | Demon Lover Diary (1980) | 1 | 4.500000 |
| 16 | 25975 | Life of Oharu, The (Saikaku ichidai onna) (1952) | 3 | 4.500000 |
| 17 | 38320 | Valerie and Her Week of Wonders (Valerie a t   den divu) (1970) | 1 | 4.500000 |
| 18 | 50477 | Testament of Orpheus, The (Testament d'Orph   ) (1960) | 1 | 4.500000 |
| 19 | 53883 | Power of Nightmares: The Rise of the Politics of Fear, The (2004) | 4 | 4.500000 |

| ranking | movieId | title | votes | stars |
|---------|---------|---|-------|----------|
| 20 | 56015 | Kansas City Confidential (1952) | 1 | 4.500000 |
| 21 | 58185 | Please Vote for Me (2007) | 1 | 4.500000 |
| 22 | 60336 | Bad Blood (Mauvais sang) (1986) | 1 | 4.500000 |
| 23 | 60990 | End of Summer, The (Kohayagawa-ke no aki) (1961) | 3 | 4.500000 |
| 24 | 61695 | Ladrones (2007) | 1 | 4.500000 |
| 25 | 63179 | Tokyo! (2008) | 1 | 4.500000 |
| 26 | 64418 | Man Named Pearl, A (2006) | 1 | 4.500000 |
| 27 | 318 | Shawshank Redemption, The (1994) | 28015 | 4.455131 |
| 28 | 32792 | Red Desert, The (Deserto rosso, Il) (1964) | 6 | 4.416667 |
| 29 | 858 | Godfather, The (1972) | 17747 | 4.415366 |
| 30 | 32657 | Man Who Planted Trees, The (Homme qui plantait des arbres, L') (1987) | 5 | 4.400000 |
| 31 | 50 | Usual Suspects, The (1995) | 21648 | 4.365854 |
| 32 | 527 | Schindler's List (1993) | 23193 | 4.363493 |
| 33 | 31309 | Rocco and His Brothers (Rocco e i suoi fratelli) (1960) | 3 | 4.333333 |
| 34 | 58808 | Last Hangman, The (2005) | 3 | 4.333333 |
| 35 | 912 | Casablanca (1942) | 11232 | 4.320424 |
| 36 | 904 | Rear Window (1954) | 7935 | 4.318651 |
| 37 | 922 | Sunset Blvd. (a.k.a. Sunset Boulevard) (1950) | 2922 | 4.315880 |
| 38 | 1212 | Third Man, The (1949) | 2967 | 4.311426 |
| 39 | 3435 | Double Indemnity (1944) | 2154 | 4.310817 |
| 40 | 1178 | Paths of Glory (1957) | 1571 | 4.308721 |
| 41 | 2019 | Seven Samurai (Shichinin no samurai) (1954) | 5190 | 4.306744 |
| 42 | 1221 | Godfather: Part II, The (1974) | 11920 | 4.301971 |
| 43 | 58559 | Dark Knight, The (2008) | 2353 | 4.297068 |
| 44 | 750 | Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb (1964) | 10627 | 4.295333 |
| 45 | 1193 | One Flew Over the Cuckoo's Nest (1975) | 13014 | 4.293261 |
| 46 | 44555 | Lives of Others, The (Das Leben der Anderen) (2006) | 1108 | 4.291065 |
| 47 | 3030 | Yojimbo (1961) | 1528 | 4.281741 |
| 48 | 1148 | Wallace & Gromit: The Wrong Trousers (1993) | 7167 | 4.275429 |
| 49 | 745 | Wallace & Gromit: A Close Shave (1995) | 5690 | 4.275308 |
| 50 | 1260 | M (1931) | 1926 | 4.274662 |

2.3 Data visualization

The most popular ratings are 4.0 and 3.0, while half stars ratings are less common.


```
# Rating histogram
y_cuts <- c(0.5, 1.0, 1.5, 2.0, 2.5)
edx %>%
  ggplot(aes(rating, fill = cut(rating, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "Rating", y = "Movies") +
```

```

ggtitle("Figure 2.1: Rating distribution per movies") +
theme(panel.grid.major = element_blank(),
      panel.grid.minor = element_blank(),
      panel.background = element_rect(fill = "white")) +
scale_x_discrete(limits = c(seq(min(edx$rating),
                                max(edx$rating),
                                max(edx$rating)/n_distinct(edx$rating))))+
scale_y_continuous(labels = paste0(y_cuts, "M"),
                   breaks = y_cuts * 10^6)

```

Figure 2.1: Rating distribution per movies



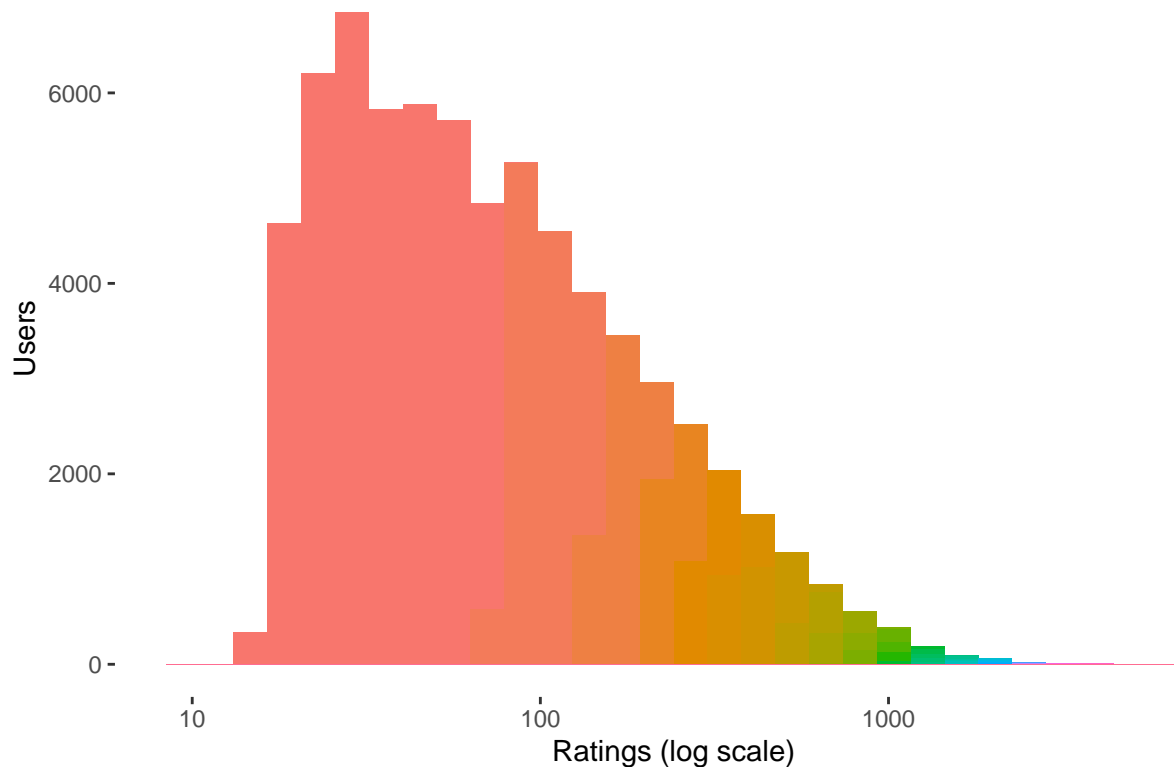
Users usually rate less than 100 different movies, and very few more than 1000.

```

# Ratings per users histogram
edx %>% count(userId) %>%
  ggplot(aes(n, fill = cut(n, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "Ratings (log scale)", y = "Users") +
  ggtitle("Figure 2.2: Ratings per users") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white")) +
  scale_x_log10()

```

Figure 2.2: Ratings per users




Mean rating is between 3.5 and 4.0.

```
# Mean rating per users histogram
edx %>% group_by(userId) %>%
  summarise(m = mean(rating)) %>%
  ggplot(aes(m, fill = cut(m, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "Mean rating", y = "Users") +
  ggtitle("Figure 2.3: Mean rating per users") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white")) +
  scale_x_discrete(limits = c(seq(min(edx$rating),
                                max(edx$rating),
                                max(edx$rating)/n_distinct(edx$rating))))
```

```
## Warning: Continuous limits supplied to discrete scale.
## Did you mean 'limits = factor(...)' or 'scale_*_continuous()'?
```



Figure 2.3: Mean rating per users



On the other hand, less than ten movies have received less than ten ratings. But also less than 400 have received more than 1000 ratings.

```
# Ratings per movies histogram
edx %>% count(movieId) %>%
  ggplot(aes(n, fill = cut(n, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "Ratings (log scale)", y = "Movies") +
  ggtitle("Figure 2.4: Ratings per movies") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white")) +
  scale_x_log10()
```


Figure 2.4: Ratings per movies



Mean rating is near 3.5.

```
# Mean rating per movies histogram
edx %>% group_by(movieId) %>%
  summarise(m = mean(rating)) %>%
  ggplot(aes(m, fill = cut(m, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "Mean rating", y = "Movies") +
  ggtitle("Figure 2.5: Mean rating per movies") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white")) +
  scale_x_discrete(limits = c(seq(min(edx$rating),
                                max(edx$rating),
                                max(edx$rating)/n_distinct(edx$rating))))
```

Figure 2.5: Mean rating per movies



2.4 Data calculations

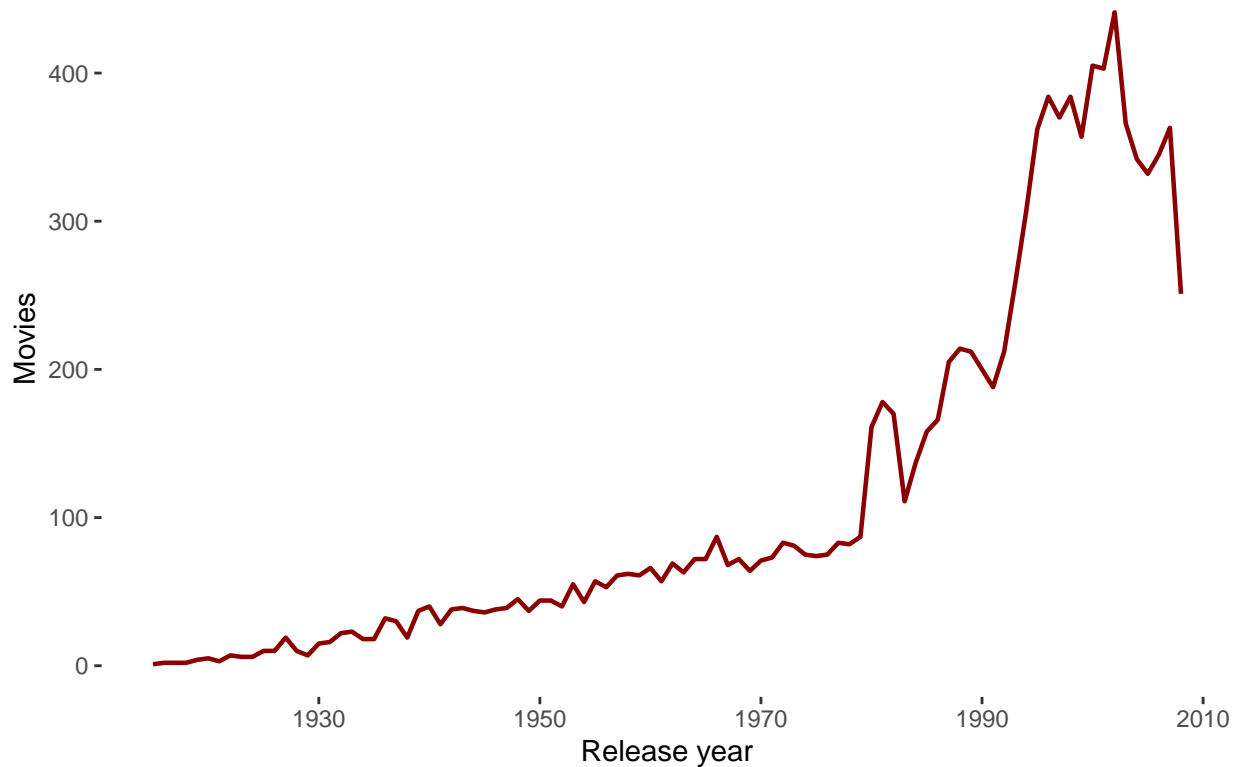
Notice that the data's timestamp can be transformed using January 1st, 1970 as initial date. We can also extract the release year from the title, and using both years we can calculate the rate timestamp.

```
# Dates addition
edx <- edx %>%
  mutate(timestamp = as.POSIXct(timestamp,
                                origin = "1970-01-01",
                                tz = "GMT"),
         year_movie = as.numeric(substr(title,
                                         nchar(title)-4,
                                         nchar(title)-1)),
         year Rated = as.numeric(format(timestamp, "%Y")),
         rate_ts = year Rated - year_movie)
```

There are less than 100 movies per year for movies released before 1980, but more than 200 for those released after 1990.

```
edx %>%
  select(movieId, year_movie) %>%
  unique() %>%
  group_by(year_movie) %>%
  summarise(count = n()) %>%
  ggplot(aes(year_movie, count)) +
  geom_line(color = "darkred",
            size = 0.8) +
  labs(x = "Release year", y = "Movies") +
  ggtitle("Figure 2.6: Release year per movies") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))
```

Figure 2.6: Release year per movies



2.5 Data modeling

Now we make a naive prediction based on the average rating, and from there, we can build the model adding movie, user, release year and rate timestamp effects. Finally, we regularize these effects to minimize the RMSE.

```
# Summary statistics
mu <- mean(edx$rating)
```

The average rating ($\mu = 3.5124652$) is used to make a naive prediction. We obtain RMSE = 1.06, meaning our predictions are on average more than one star far from the actual value.

```
# Average prediction
rmse_naive <- RMSE(edx$rating, mu)
rmse_results <- tibble(Model = "01. Average only",
                       RMSE = rmse_naive,
                       Goal = ifelse(rmse_naive < 0.8649, "Under", "Over"))
```

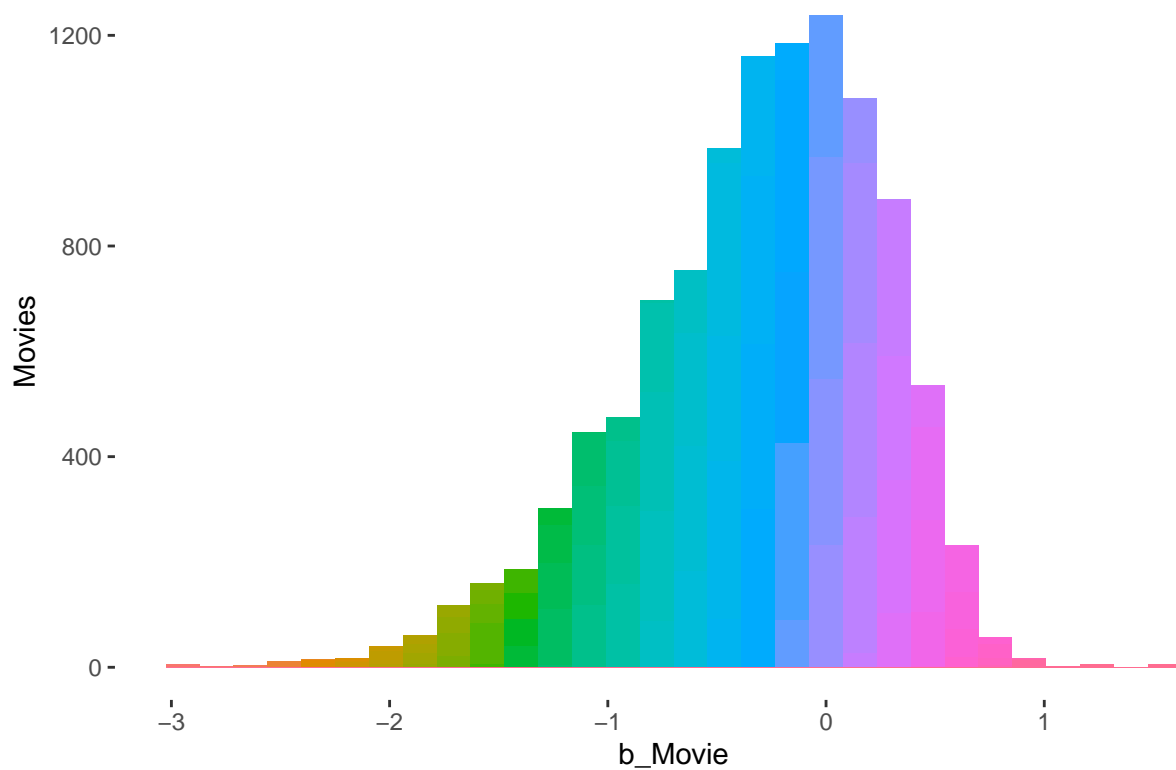
| Model | RMSE | Goal |
|------------------|----------|------|
| 01. Average only | 1.060331 | Over |

Now we will start incorporating the movie effect, which generates a RMSE = 0.9423.

```
# Movie effect prediction
movie_avgs <- edx %>%
  group_by(movieId) %>%
  summarise(b_Movie = mean(rating - mu))

movie_avgs %>%
  ggplot(aes(b_Movie, fill = cut(b_Movie, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "b_Movie", y = "Movies") +
  ggtitle("Figure 3.1: Movies per computed b_Movie") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))
```

Figure 3.1: Movies per computed b_Movie



```
predicts_movie <- mu + edx %>%
  left_join(movie_avgs, by = "movieId") %>%
  pull(b_Movie)

rmse_model1 <- RMSE(predicts_movie, edx$rating)

rmse_results <- rmse_results %>%
  bind_rows(tibble(Model = "02. Movie Effect",
                    RMSE = rmse_model1,
                    Goal = ifelse(rmse_model1 < 0.8649, "Under", "Over")))
```

| Model | RMSE | Goal |
|------------------|-----------|------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |

Then we add the user effect. With this two, we get a RMSE = 0.8567, which is lower than the target.

```
# Movie and user effect prediction
user_avgs <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
```

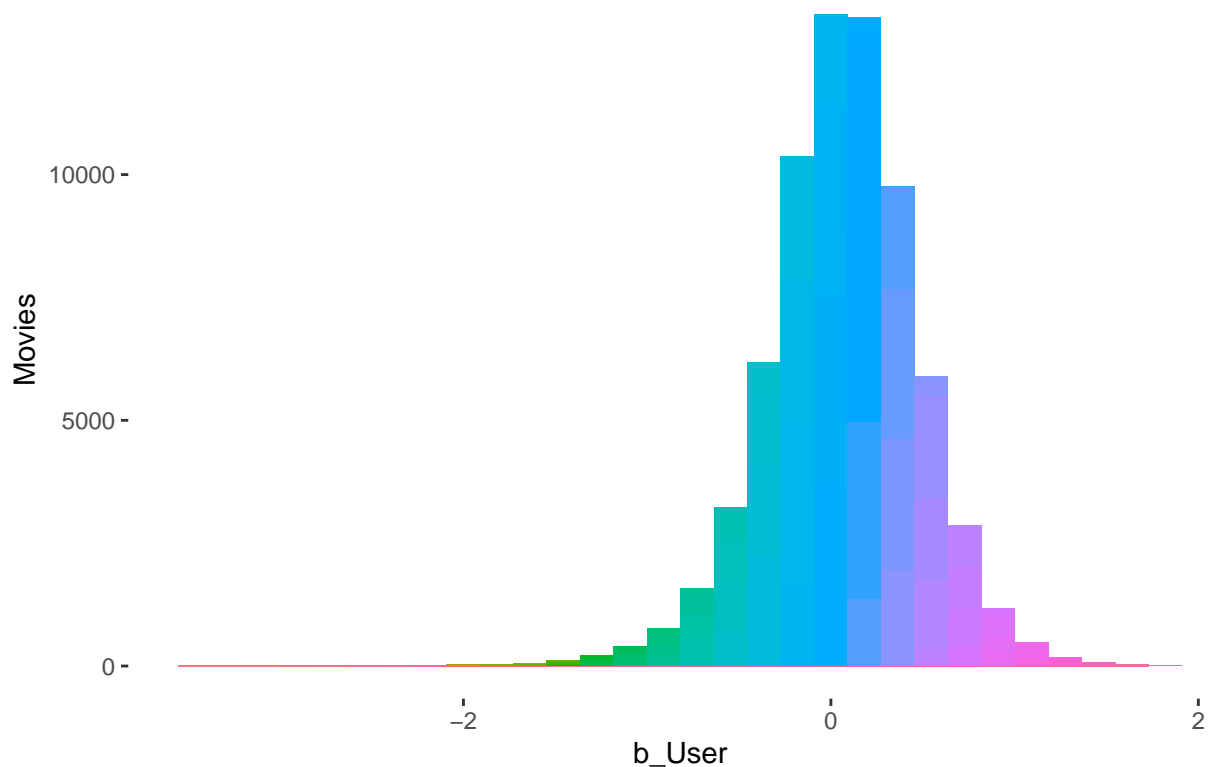
```

group_by(userId) %>%
  summarise(b_User = mean(rating - b_Movie - mu))

user_avgs %>%
  ggplot(aes(b_User, fill = cut(b_User, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "b_User", y = "Movies") +
  ggtitle("Figure 3.2: Movies per computed b_User") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))

```

Figure 3.2: Movies per computed b_User



```

predicts_movie_user <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
  left_join(user_avgs, by = "userId") %>%
  mutate(pred = mu + b_Movie + b_User) %>%
  pull(pred)

rmse_model2 <- RMSE(predicts_movie_user,
                    edx$rating)
rmse_results <- rmse_results %>%

```

```
bind_rows(tibble(Model = "03. +User Effect",
                  RMSE = rmse_model2,
                  Goal = ifelse(rmse_model2 < 0.8649, "Under", "Over")))
```

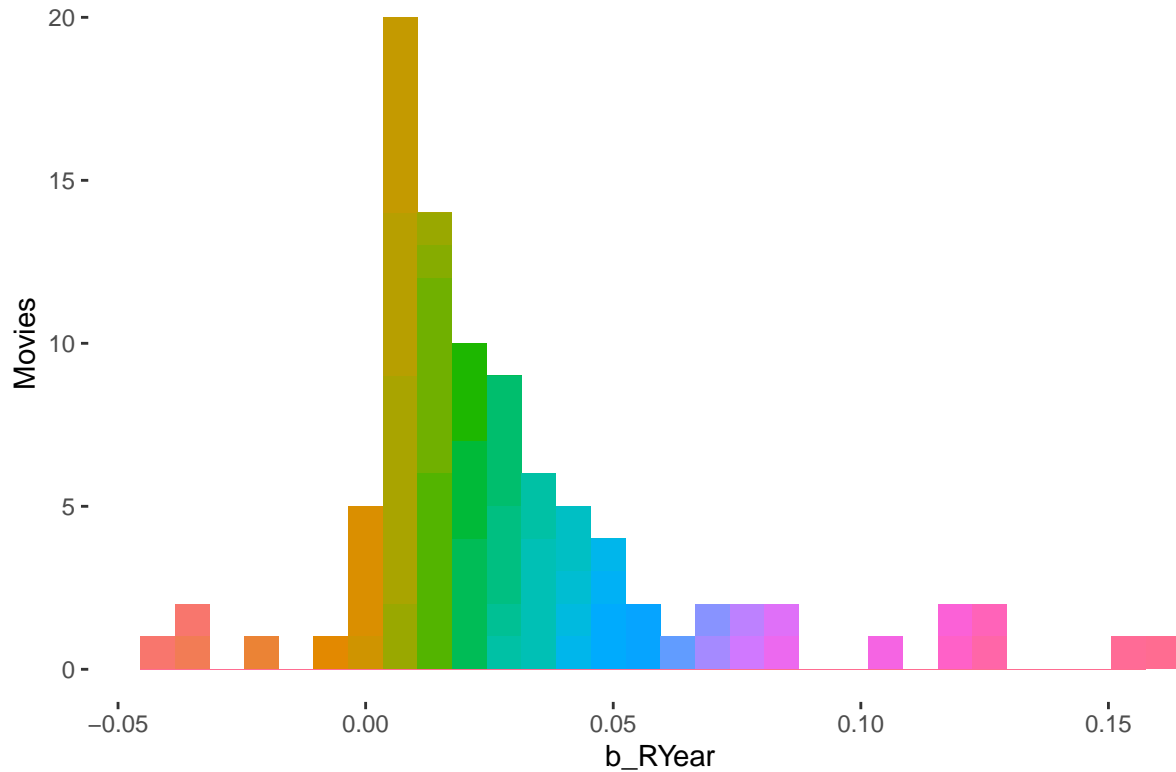
| Model | RMSE | Goal |
|------------------|-----------|-------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |
| 03. +User Effect | 0.8567039 | Under |

Then, we incorporate the release year effect, getting a RMSE = 0.8564. We can see how β *RYear* values are approaching 0.

```
# Movie, user and release year effect prediction
RYear_avgs <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
  left_join(user_avgs, by = "userId") %>%
  group_by(year_movie) %>%
  summarise(b_RYear = mean(rating - b_User - b_Movie - mu))

RYear_avgs %>%
  ggplot(aes(b_RYear, fill = cut(b_RYear, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "b_RYear", y = "Movies") +
  ggtitle("Figure 3.3: Movies per computed b_RYear") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))
```


Figure 3.3: Movies per computed b_RYear



```
predicts_movie_user_RYear <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
  left_join(user_avgs, by = "userId") %>%
  left_join(RYear_avgs, by = "year_movie") %>%
  mutate(pred = mu + b_Movie + b_User + b_RYear) %>%
  pull(pred)

rmse_model3 <- RMSE(predicts_movie_user_RYear,
  edx$rating)
rmse_results <- rmse_results %>%
  bind_rows(tibble(Model = "04. +RYear Effect",
    RMSE = rmse_model3,
    Goal = ifelse(rmse_model3 < 0.8649, "Under", "Over")))
```

| Model | RMSE | Goal |
|-------------------|-----------|-------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |
| 03. +User Effect | 0.8567039 | Under |
| 04. +RYear Effect | 0.8563777 | Under |


Our fourth and final effect is rate timestamp, giving a $RMSE = 0.8561$. Most values of β_{rateTS}

are closer to 0 from less than 0.1.

```
# Movie, user, release year and rate timestamp effect prediction
rateTS_avgs <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
  left_join(user_avgs, by = "userId") %>%
  left_join(RYear_avgs, by = "year_movie") %>%
  group_by(rate_ts) %>%
  summarise(b_rateTS = mean(rating - b_RYear - b_User - b_Movie - mu))

rateTS_avgs %>%
  ggplot(aes(b_rateTS, fill = cut(b_rateTS, 100))) +
  geom_histogram(bins = 30, show.legend = F) +
  theme_update() +
  labs(x = "b_rateTS", y = "Movies") +
  ggtitle("Figure 3.4: Movies per computed b_rateTS") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))
```

Figure 3.4: Movies per computed b_rateTS



```
predicts_movie_user_RYear_rateTS <- edx %>%
  left_join(movie_avgs, by = "movieId") %>%
```

```

left_join(user_avgs, by = "userId") %>%
left_join(RYear_avgs, by = "year_movie") %>%
left_join(rateTS_avgs, by = "rate_ts") %>%
mutate(pred = mu + b_rateTS + b_RYear + b_User + b_Movie) %>%
pull(pred)

rmse_model4 <- RMSE(predicts_movie_user_RYear_rateTS,
                    edx$rating)
rmse_results <- rmse_results %>%
  bind_rows(tibble(Model = "05. +RateTS Effect",
                    RMSE = rmse_model4,
                    Goal = ifelse(rmse_model4 < 0.8649, "Under", "Over"))))

```

| Model | RMSE | Goal |
|--------------------|-----------|-------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |
| 03. +User Effect | 0.8567039 | Under |
| 04. +RYear Effect | 0.8563777 | Under |
| 05. +RateTS Effect | 0.8560683 | Under |

Now we can regularize our model. We use cross validation with a vector of tuning parameters from 0 to 5, increasing on 0.25. The result is an optimal $\lambda = 0.5$.

```

# Regularized effects
lambdas <- seq(0, 5, 0.25)

RMSEs <- sapply(lambdas, function(l){

  b_Movie <- edx %>%
    group_by(movieId) %>%
    summarise(b_Movie = sum(rating - mu)/(n() + 1))

  b_User <- edx %>%
    left_join(b_Movie, by = "movieId") %>%
    group_by(userId) %>%
    summarise(b_User = sum(rating - mu - b_Movie)/(n() + 1))

  b_RYear <- edx %>%
    left_join(b_Movie, by = "movieId") %>%
    left_join(b_User, by = "userId") %>%
    group_by(year_movie) %>%
    summarise(b_RYear = sum(rating - mu - b_Movie - b_User)/(n() + 1))

  b_rateTS <- edx %>%
    left_join(b_Movie, by = "movieId") %>%

```

```

    left_join(b_User, by = "userId") %>%
    left_join(b_RYear, by = "year_movie") %>%
    group_by(rate_ts) %>%
    summarise(b_rateTS = sum(rating - mu - b_Movie - b_User - b_RYear)/(n() + 1))


predicts_ratings <- edx %>%
  left_join(b_Movie, by = "movieId") %>%
  left_join(b_User, by = "userId") %>%
  left_join(b_RYear, by = "year_movie") %>%
  left_join(b_rateTS, by = "rate_ts") %>%
  mutate(pred = mu + b_Movie + b_User + b_RYear + b_rateTS) %>%
  pull(pred)

return(RMSE(predicts_ratings, edx$rating))
})

tibble(lambdas) %>%
  bind_cols(tibble(RMSEs)) %>%
  ggplot(aes(lambdas, RMSEs)) +
  geom_point() +
  theme_update() +
  labs(x = "lambdas", y = "RMSEs") +
  ggtitle("Figure 3.5: RMSE per lambda") +
  theme(panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        panel.background = element_rect(fill = "white"))

```

Figure 3.5: RMSE per lambda



```
rmse_model5 <- min(RMSEs)
lambda <- lambdas[which.min(RMSEs)]
```

```
## [1] 0.5
```

Implementing $\lambda = 0.5$ generates a RMSE = 0.8560611. We are now 0.0088389 under our target.

```
rmse_results <- rmse_results %>%
  bind_rows(tibble(Model = "06. Regularized",
    RMSE = rmse_model5,
    Goal = ifelse(rmse_model5 < 0.8649, "Under", "Over")))
```

| Model | RMSE | Goal |
|--------------------|-----------|-------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |
| 03. +User Effect | 0.8567039 | Under |
| 04. +RYear Effect | 0.8563777 | Under |
| 05. +RateTS Effect | 0.8560683 | Under |
| 06. Regularized | 0.8560611 | Under |

3 Results

A recommendation system “is a subclass of information filtering system that seeks to predict the ‘rating’ or ‘preference’ a user would give to an item” (Shi, 2020). In order to make this prediction, we have built an ML regularized model that considers four different effects: Movie, User, Release Year & Rate Timestamp.

Until this moment, we have only used our edx set, obtaining an RMSE of 0.8560611. In this section, we implement the model on the Validation Set.

We start replicating the effects while incorporating the regularized λ .

```
# Replication with optimal lambda
b_Movie <- edx %>%
  group_by(movieId) %>%
  summarise(b_Movie = sum(rating - mu)/(n() + lambda))

b_User <- edx %>%
  left_join(b_Movie, by = "movieId") %>%
  group_by(userId) %>%
  summarise(b_User = sum(rating - mu - b_Movie)/(n() + lambda))

b_RYear <- edx %>%
  left_join(b_Movie, by = "movieId") %>%
  left_join(b_User, by = "userId") %>%
  group_by(year_movie) %>%
  summarise(b_RYear = sum(rating - mu - b_Movie - b_User)/(n() + lambda))

b_RateTS <- edx %>%
  left_join(b_Movie, by = "movieId") %>%
  left_join(b_User, by = "userId") %>%
  left_join(b_RYear, by = "year_movie") %>%
  group_by(rate_ts) %>%
  summarise(b_rateTS = sum(rating - mu - b_Movie - b_User - b_RYear)/(n() + lambda))
```

Then we calculate the time variables in the validation set.

```
# Validation data arrangement
validation <- validation %>%
  mutate(timestamp = as.POSIXct(timestamp,
                                origin = "1970-01-01",
                                tz = "GMT"),
         year_movie = as.numeric(substr(title,
                                         nchar(title)-4,
                                         nchar(title)-1)),
         year Rated = as.numeric(format(timestamp, "%Y")),
         rate_ts = year Rated - year_movie)
```

Finally we apply the model to the validation set.

```
# Final prediction
predicts_validation <- validation %>%
  left_join(movie_avgs, by = "movieId") %>%
  left_join(user_avgs, by = "userId") %>%
  left_join(RYear_avgs, by = "year_movie") %>%
  left_join(rateTS_avgs, by = "rate_ts") %>%
  mutate(pred = mu + b_rateTS + b_RYear + b_User + b_Movie) %>%
  pull(pred)

rmse_model_final <- RMSE(predicts_validation,
  validation$rating)
```

We get an RMSE = 0.8647026. Because our goal was to generate a model with RMSE < 0.8649, we reach an RMSE 0.0001974422 under our target.

```
rmse_results <- rmse_results %>%
  bind_rows(tibble(Model = "07. Validation set",
    RMSE = rmse_model_final,
    Goal = ifelse(rmse_model_final < 0.8649, "Under", "Over")))
```

| Model | RMSE | Goal |
|--------------------|-----------|-------|
| 01. Average only | 1.0603313 | Over |
| 02. Movie Effect | 0.9423475 | Over |
| 03. +User Effect | 0.8567039 | Under |
| 04. +RYear Effect | 0.8563777 | Under |
| 05. +RateTS Effect | 0.8560683 | Under |
| 06. Regularized | 0.8560611 | Under |
| 07. Validation set | 0.8647026 | Under |

4 Conclusion

Across this project, we first described the MovieLens 10M dataset while presenting the four variables (two existing, two calculated) we focused on. Then, we did some exploratory data analysis gaining insights for the model. Finally, we developed a recommendation system that considers the effect depending on the movie, the user, the release year, and the rate timestamp. We regularized the effects with a tuning parameter of 0.5, and got a RMSE = 0.8647026 once applied to the Validation Set.

Even if the model reached its goal with a 1.9744224×10^{-4} margin, it still has plenty of opportunity areas to improve on future work:

1. As seen, many top places are occupied by movies with less than five ratings, making the model to consider movie characteristics that are less ranked.

2. The genres variable isn't used in this model. Desagregating this column could give more insight for better predictions.
3. Other option is to work on data enrichment, developing APIs to get more characteristics about the movies.
4. GroupLens is constantly improving their software, so downloading bigger and more recent databases could give more insights for predictions.
5. Meng (2020) and other bloggers constantly publish more advanced ML techniques that are able to refine the analysis.

5 References

- Harper, F.M. & Konstan, J.A. (2015). "The MovieLens Datasets: History and Context". ACM Transactions on Interactive Intelligent Systems: University of Minnesota. Available in: <https://dl.acm.org/doi/10.1145/2827872>
- Irizarry, R. A. (2020). "Introduction to Data Science: Data Analysis and Prediction Algorithms with R". Available in: <https://rafalab.github.io/dsbook/>
- Meng, M. (2020). "Movie Recommendation System based on MovieLens: Leveraged Natural Language Processing (NLP) and Computer Vision (CV)". Towards Data Science. Available in: <https://towardsdatascience.com/movie-recommendation-system-based-on-movielens-efodf58ocdoe>
- Shi, W. (2020). "Recommendation Systems: A Review". Towards Data Science. Available in: <https://towardsdatascience.com/recommendation-systems-a-review-d4592b6caf4b>

6 Appendix

```
sessionInfo()
```

```
## R version 4.0.3 (2020-10-10)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
## Running under: Windows 10 x64 (build 19041)
##
## Matrix products: default
##
## Random number generation:
##  RNG:      Mersenne-Twister
##  Normal:   Inversion
##  Sample:   Rounding
##
## locale:
## [1] LC_COLLATE=English_United States.1252
```



```

## [2] LC_CTYPE=English_United States.1252
## [3] LC_MONETARY=English_United States.1252
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United States.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] knitr_1.30      data.table_1.13.2  caret_6.0-86      lattice_0.20-41
## [5] forcats_0.5.0   stringr_1.4.0      dplyr_1.0.2       purrr_0.3.4
## [9] readr_1.4.0     tidyr_1.1.2        tibble_3.0.4      ggplot2_3.3.2
## [13] tidyverse_1.3.0
##
## loaded via a namespace (and not attached):
## [1] httr_1.4.2      jsonlite_1.7.1    splines_4.0.3
## [4] foreach_1.5.1   prodlim_2019.11.13 modelr_0.1.8
## [7] assertthat_0.2.1 highr_0.8          stats4_4.0.3
## [10] cellranger_1.1.0 yaml_2.2.1         ipred_0.9-9
## [13] pillar_1.4.7    backports_1.2.0   glue_1.4.2
## [16] pROC_1.16.2     digest_0.6.27     rvest_0.3.6
## [19] colorspace_2.0-0 recipes_0.1.15     htmltools_0.5.0
## [22] Matrix_1.2-18   plyr_1.8.6         timeDate_3043.102
## [25] pkgconfig_2.0.3 broom_0.7.2        haven_2.3.1
## [28] scales_1.1.1    gower_0.2.2       lava_1.6.8.1
## [31] farver_2.0.3    generics_0.1.0     ellipsis_0.3.1
## [34] withr_2.3.0     nnet_7.3-14        cli_2.2.0
## [37] survival_3.2-7  magrittr_2.0.1     crayon_1.3.4
## [40] readxl_1.3.1    evaluate_0.14      fs_1.5.0
## [43] fansi_0.4.1     nlme_3.1-150       MASS_7.3-53
## [46] xml2_1.3.2      class_7.3-17       tools_4.0.3
## [49] hms_0.5.3       lifecycle_0.2.0    munsell_0.5.0
## [52] reprex_0.3.0    compiler_4.0.3     rlang_0.4.9
## [55] grid_4.0.3      iterators_1.0.13    rstudioapi_0.13
## [58] labeling_0.4.2  rmarkdown_2.5      gtable_0.3.0
## [61] ModelMetrics_1.2.2.2 codetools_0.2-18   DBI_1.1.0
## [64] reshape2_1.4.4  R6_2.5.0           lubridate_1.7.9.2
## [67] utf8_1.1.4      stringi_1.5.3      Rcpp_1.0.5
## [70] vctrs_0.3.5     rpart_4.1-15       dbplyr_2.0.0
## [73] tidyselect_1.1.0 xfun_0.19

```